

Programming Paradigms CSI2120 – Winter 2018

**Jochen Lang
EECS, University of Ottawa
Canada**

Université d'Ottawa | University of Ottawa



uOttawa

L'Université canadienne
Canada's university



uOttawa.ca

System Programming: Go

- Concurrency and Parallelism
- Goroutines
- Channels

Concurrent Programming

- **An application is a process running on a machine**
 - A process is an independently executing entity that runs in its own address space.
- **A process is composed of one or more operating system threads**
 - Threads are concurrently executing entities that share the same address space.

Execution Thread

- **An execution thread is a sequence of executable statements and it may or may not interact with other threads**
 - Threads often share some variables (or resources) with other threads
 - Threads often (but not necessarily) have a limited life-time within a program execution
- **A thread can be blocked:**
 - Accessing a shared variable or resource used currently by another thread
 - Waiting for the result or completion of another thread
- **An application program can often be divided into a process and potentially many threads**

Parallel Programming vs. Concurrent Programming

- **Concurrent programming means expressing program structure such that is organized into independently executing actions**
 - Concurrent programming can also target a single processor
 - The processes or threads then run in turn over time according to a schedule
- **If processes or threads are running on different processors or cores simultaneously, we have a parallel program**
- **If a program is executed on multiple machines with loose interaction, it is commonly called distributed programming**
- **If a program is executed on a graphics card or a tightly integrated cluster, it is often called massively parallel programming**

Concurrent Programming Languages

- **A concurrent programming language must support:**
 - the creation and execution of processes and threads
 - synchronization operations
 - Cooperative synchronization: A process waits for the execution of another before continuing its own execution.
 - deterministic concurrency
 - Competitive synchronization: Multiple processes use the same resource with some form of locking mechanism for mutual exclusion.
 - non-deterministic concurrency
 - data communication between threads and processes
 - may use mechanisms using inter-process communication defined by the operating system

Reminder: Threads in Java

- Option implementing interface runnable

```
public class HelloRunnable
    implements Runnable {

    public void run() {
        System.out.println(
            "Running in a
            thread");
    }

    public static void
        main( String args[]) {
        (new Thread( new
            HelloRunnable())) .start();
    }
}
```

- Option subclassing Thread

```
public class HelloThread
    extends Thread {

    public void run() {
        System.out.println(
            "Running in a
            thread!");
    }

    public static void
        main(String args[]) {
        (new
            HelloThread()) .start();
    }
}
```

Level of Concurrency

- **At the (set of) statement level:**
 - Sets of statements are executed concurrently while the main process suspends its execution. All threads share the same data (OpenMP)
- **At the sub-program level:**
 - New process is created for running a subroutine. Once the new process starts, the calling process continues its execution. Requires a synchronization mechanism.
- **At the object level:**
 - Each object instance of a class competes for resources and methods on different objects run concurrently. Class variables (attributes) are not shared.
- **At the program level:**
 - Parent process run one or more child processes. Child process Id must be known by parent. Data may be shared.

Type of Concurrency

- **Physical**
 - Multiple processes/threads share different processors or cores
- **Logical**
 - Multiple processes/threads share execution time on a single processor.
- **Distributed**
 - Processes/threads of an application share several machines over a network

Concurrency in Go

- **Two mechanisms are provided**
 - Non-deterministic Concurrency
 - More traditional threads with low-level synchronization
 - Mutex in synch package
 - Use is discouraged in Go
 - Deterministic Concurrency
 - Communicating Sequential Processes (CSP)
 - Based on message passing between threads
 - goroutines and channels
 - Recommended approach

Concurrency in Go

- **CSP**

- Based on the idea that avoiding data sharing will avoid the biggest problem in concurrent programming
- Threads in CSP do not communicate by sharing data
- Rather they share data by communicating
- Timing of threads is based on messaging between threads

Goroutines

- **Parts of a Go program that run concurrently are organized in *goroutines***
 - *goroutines* can run in several threads
 - several *goroutines* can run in one thread
 - no 1:1 correspondence to OS threads
- ***goroutines* run in the same address space**
 - shared memory access can be used but is discouraged
- ***goroutines* are designed to be light-weight**
 - inexpensive to create
 - automatic (segmented) stack management
- **A *goroutine* can be implemented as a function or method**
- **A *goroutine* is invoked using the keyword `go`**

Calling a *Goroutine*

- ***Goroutines* are functions or methods called with `go`**
 - The default number of OS threads is one for all goroutines
 - Can be changed with the environment variable `GOMAXPROCS`
 - Can also be changed with the runtime package

```
import "runtime"

func main() {
    // change max number of OS threads to 3
    runtime.GOMAXPROCS(3)
    ...
    go foo()
    ...
}
```

Example: Calling *goroutines*

```
func main() {  
    runtime.GOMAXPROCS(3)  
    sTime := time.Now(); // time it  
    fmt.Println("Starting")  
    go letters() // goroutine A  
    go numbers() // goroutine B  
    fmt.Println("Waiting ...")  
    time.Sleep(2*time.Second)  
    fmt.Println("\nDone\n")  
    eTime := time.Now();  
    fmt.Printf("Run time: %s", eTime.Sub(sTime))  
}
```

Example: *goroutines*

```
func numbers() {  
    for number := 1; number < 27; number++ {  
        // pause before every print  
        time.Sleep(10*time.Millisecond)  
        fmt.Printf("%d ", number)  
    }  
}
```

```
func letters() {  
    for char := 'a'; char < 'a'+26; char++ {  
        time.Sleep(10)  
        fmt.Printf("%c ", char)  
    }  
}
```

Example Execution

Starting

Waiting ...

```
a b c d e f g h i j l k l m n o p q r s t 2 u v w
x y z 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19
20 21 22 23 24 25 26
```

Done

Run time: 2s

- **Running it without *goroutines* (sequential)**

Starting

```
a b c d e f g h i j k l m n o p q r s t u v w x y
z 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19
20 21 22 23 24 25 26 Waiting ...
```

Done

Run time: 2.286s

Communication between *goroutines*

- *goroutines* are designed to effectively communicate by message passing for
 - the exchange of information
 - for synchronizing the execution
- **Implemented in Go through channels**
 - channels are similar to pipes

Concept of *channels*

- **Data is passed around through channels**
 - only one *goroutine* has access to a channel at any given time
 - communication is synchronized by design
 - no race conditions can occur
- **Ownership of the data is passed around as well**
- **A channel is a data queue (FIFO)**
- **Channels are typed (only one data type can be passed around on a channel)**

Channel Declaration

- **Declaring a channel only creates a reference**

- Use make to allocate space for the channel

- Example: Channel for strings

```
var ch chan string
ch = make(chan string)
```

- Or shorter with initializer declaration

```
ch := make(chan string)
```

- By default, a channel has a capacity of 1 element

- Channels are first-class object

- To send or receive, use the arrow operator

```
ch <- str // send
str = <- ch // receive
```

Example: Communicating *goroutines* through channels

```
func main() {
    ch := make(chan string)
    go sendString(ch)
    go recieveString(ch)
    time.Sleep(1*time.Second)
}

func sendString(ch chan string) {
    ch <- "Ottawa"
    ch <- "Toronto"
    ch <- "Gatineau"
    ch <- "Casselman"
}

func recieveString (ch chan string) {
    var str string
    for {
        str= <-ch
        fmt.Printf("%s ", str)
    }
}
```

Range Loop applied to a Channel

- We can loop over the incoming elements
- Loops over the channel until it is closed
- Example uses a lambda as go routine

```
func recieveString(ch chan string) {  
    go func() {  
        for str := range ch {  
            fmt.Printf("%s ", str)  
        }  
    }()  
}
```

Closing channels

- **Channels may be explicitly closed**

```
func sendString(strArr []string) chan string {  
    ch := make(chan string)  
    go func() { // start a lambda in a go routine  
        for _, s := range strArr {  
            ch <- s  
        }  
        close(ch)  
    }()  
    return ch  
}
```

- **We can test if the channel has been closed**

```
for {  
    str, ok := <- ch  
    if !ok {  
        break;  
    }  
    fmt.Printf("%s ", str)  
}
```

Synchronization across Channels

- We can use **select** for non-blocking channel i/o. Similar to a switch but depending on which channel receives or ready to send an element

```
Forever:
    for {
        select {
            case str, ok := <- ch1:
                if !ok {
                    break Forever
                }
                fmt.Printf("%s \n", str )
            case str, ok := <- ch2:
                if !ok {
                    break Forever
                }
                fmt.Printf("%s \n", str )
        }
    }
}
```

Timers in Go

- **A timer in Go has a channel**
- **We can check its status or read from its channel**

```
    outOfHere := time.NewTimer(time.Second * 2)
Forever:
    for {
        select {
            ... \\ Omitted
            case <-outOfHere.C:
                fmt.Printf("Time is up\n" )
                break Forever
        }
    }
```


Summary

- **Concurrency and Parallelism**
- **Goroutines**
- **Channels**
- **Channel Synchronization**
- **Timer**